

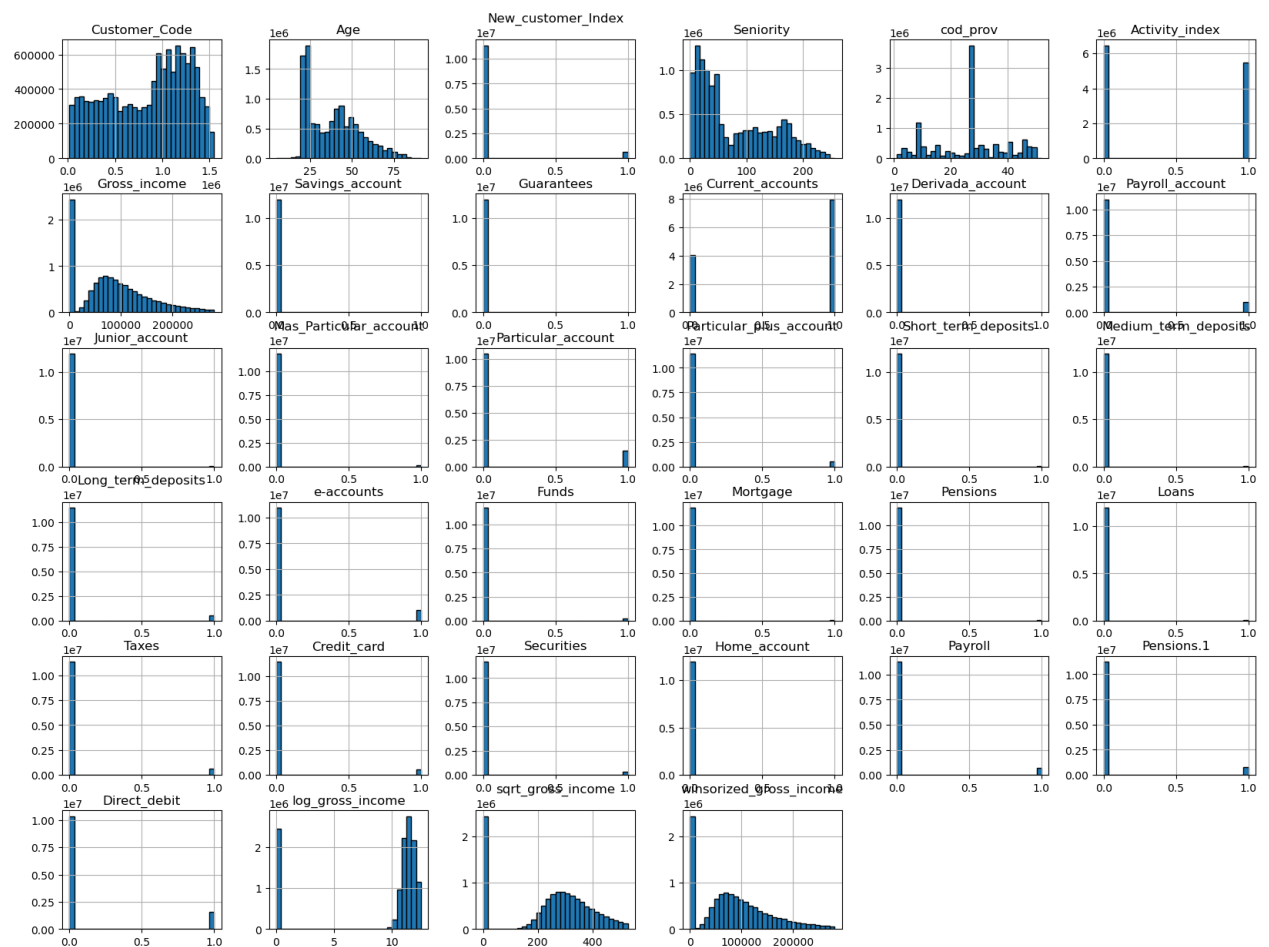
```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
import matplotlib.pyplot as plt
```

```
In [2]: # Read the CSV file
df_train = pd.read_csv("df_no_outliers.csv")
```

```
In [3]: # 2 - Univariate Analysis
# Distribution of numerical features
df_train.hist(bins=30, figsize=(20, 15), edgecolor='black')
plt.suptitle('Distribution of Numerical Features')
plt.show()

# Summary statistics for categorical columns
print("Categorical Columns Description:")
print(df_train.describe(include=['category']))
```

Distribution of Numerical Features



Categorical Columns Description:

```
-----
ValueError
)
```

Cell In[3], line 9

Traceback (most recent call last)

```

7 # Summary statistics for categorical columns
8 print("Categorical Columns Description:")
----> 9 print(df_train.describe(include=['category']))

```

File /opt/anaconda3/lib/python3.11/site-packages/pandas/core/generic.py:11552, in NDFrame.describe(self, percentiles, include, exclude)

```

11310 @final
11311 def describe(
11312     self,
11313     (...)
11315     exclude=None,
11316 ) -> Self:
11317     """
11318     Generate descriptive statistics.
11319     (...)
11550     max          NaN          3.0
11551     """
> 11552     return describe_ndframe(
11553         obj=self,
11554         include=include,
11555         exclude=exclude,
11556         percentiles=percentiles,
11557     ).__finalize__(self, method="describe")

```

File /opt/anaconda3/lib/python3.11/site-packages/pandas/core/methods/describe.py:97, in describe_ndframe(obj, include, exclude, percentiles)

```

90 else:
91     describer = DataFrameDescriber(
92         obj=cast("DataFrame", obj),
93         include=include,
94         exclude=exclude,
95     )
----> 97 result = describer.describe(percentiles=percentiles)
98 return cast(NDFrameT, result)

```

File /opt/anaconda3/lib/python3.11/site-packages/pandas/core/methods/describe.py:173, in DataFrameDescriber.describe(self, percentiles)

```

170     ldesc.append(describe_func(series, percentiles))
172 col_names = reorder_columns(ldesc)
--> 173 d = concat(
174     [x.reindex(col_names, copy=False) for x in ldesc],
175     axis=1,
176     sort=False,
177 )
178 d.columns = data.columns.copy()
179 return d

```

File /opt/anaconda3/lib/python3.11/site-packages/pandas/core/reshape/concat.py:380, in concat(objs, axis, join, ignore_index, keys, levels, names, verify_integrity, sort, copy)

```

377 elif copy and using_copy_on_write():
378     copy = False
--> 380 op = _Concatenator(
381     objs,
382     axis=axis,
383     ignore_index=ignore_index,

```

```

384     join=join,
385     keys=keys,
386     levels=levels,
387     names=names,
388     verify_integrity=verify_integrity,
389     copy=copy,
390     sort=sort,
391 )
393 return op.get_result()

```

File /opt/anaconda3/lib/python3.11/site-packages/pandas/core/reshape/concat.py:443, in _Concatenator.__init__(self, objs, axis, join, keys, levels, names, ignore_index, verify_integrity, copy, sort)

```

440 self.verify_integrity = verify_integrity
441 self.copy = copy
--> 443 objs, keys = self._clean_keys_and_objs(objs, keys)
445 # figure out what our result ndim is going to be
446 ndims = self._get_ndims(objs)

```

File /opt/anaconda3/lib/python3.11/site-packages/pandas/core/reshape/concat.py:505, in _Concatenator._clean_keys_and_objs(self, objs, keys)

```

502     objs_list = list(objs)
504 if len(objs_list) == 0:
--> 505     raise ValueError("No objects to concatenate")
507 if keys is None:
508     objs_list = list(com.not_none(*objs_list))

```

ValueError: No objects to concatenate

```

In [ ]: # Check if 'Gross_income' column exists
if 'Gross_income' in df_train.columns:
    # Apply logarithmic transformation to the 'Gross_income' column
    log_data = np.log1p(df_train['Gross_income'])


    # Create a histogram for the log-transformed 'Gross_income' column
    plt.figure(figsize=(10, 6))
    sns.histplot(log_data, kde=True)
    plt.title('Log-Transformed Distribution of Gross Income')
    plt.xlabel('Log of Gross Income')
    plt.ylabel('Frequency')
    plt.ylim(0, 500000) # Set the y-axis limit
    plt.show()
else:
    print("The 'Gross_income' column does not exist in the dataset.")

```

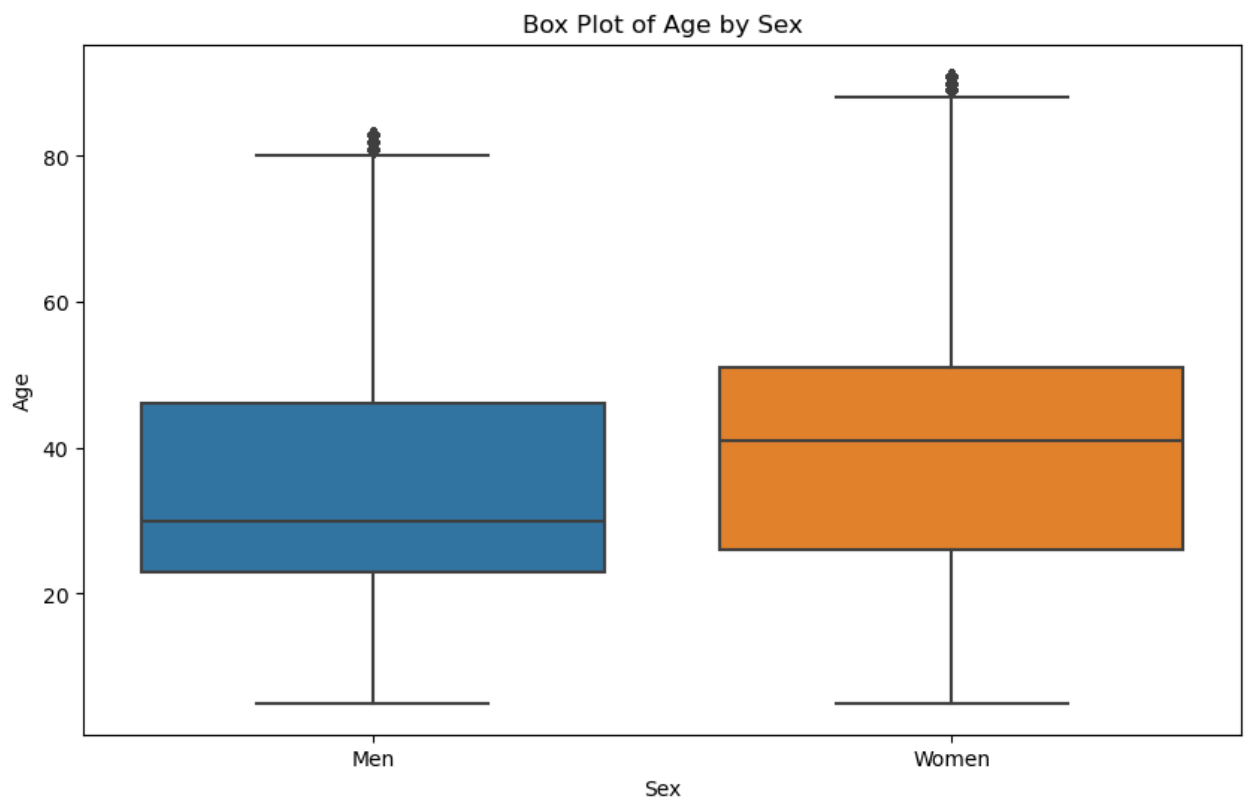
```

In [20]: # Check if Seniority column exists
if 'Seniority' in df_train.columns:
    # Create a histogram for the Seniority column
    plt.figure(figsize=(10, 6))
    sns.histplot(df_train['Seniority'], kde=True)
    plt.title('Distribution of Seniority')
    plt.xlabel('Seniority')
    plt.ylabel('Frequency')
    plt.show()
else:
    print("The Seniority column does not exist in the dataset.")

```



```
In [21]: # Plot the cleaned data
plt.figure(figsize=(10, 6))
sns.boxplot(x='Sex', y='Age', data=df_train)
plt.title('Box Plot of Age by Sex')
plt.xlabel('Sex')
plt.ylabel('Age')
plt.show()
```




```
In [22]: # Calculate the count of each channel
channel_counts = df_train['Channel_used_to_join'].value_counts()
```

```
# Get the top 8 channels
top_8_channels = channel_counts.head(8).index

# Filter the DataFrame to include only the top 5 channels
df_top_8_channels = df_train[df_train['Channel_used_to_join'].isin(top_8_

# Create the box plot
plt.figure(figsize=(14, 7))
sns.boxplot(x='Channel_used_to_join', y='Age', data=df_top_8_channels)
plt.title('Box Plot of Age by Top 8 Channels Used to Join')
plt.xlabel('Channel Used to Join')
plt.ylabel('Age')
plt.show()
```



```
In [23]: # List of columns to exclude from the correlation matrix
columns_to_exclude = ['Customer_Code'] # Add other columns to exclude if

# Drop the specified columns and select numerical columns
relevant_numerical_df = df_train.drop(columns=columns_to_exclude).select_


# Compute the correlation matrix
corr = relevant_numerical_df.corr()

# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(corr, dtype=bool))

# Set up the matplotlib figure
plt.figure(figsize=(15, 10))

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, annot=True, fmt='.2f', cmap='coolwarm', vmin=
            square=True, linewidths=.5, cbar_kws={"shrink": .5})

plt.title('Correlation Matrix')
plt.show()
```



```
In [24]: # Display the first few rows of the dataframe
print("First few rows of the dataset:")
print(df_train.head())

# Display basic information about the dataframe
print("\nBasic Information:")
print(df_train.info())

# Display summary statistics
print("\nSummary Statistics:")
print(df_train.describe())

# Check for missing values
print("\nMissing Values:")
print(df_train.isnull().sum())
```

First few rows of the dataset:

	fecha_data	Customer_Code	Employee_index	Country_of_Residence	Sex	Age
0	2015-01-28	851959	Not employed	ES	Men	36
1	2015-02-28	851959	Not employed	ES	Men	36
2	2015-03-28	851959	Not employed	ES	Men	36
3	2015-04-28	851959	Not employed	ES	Men	36
4	2015-05-28	851959	Not employed	ES	Men	36

	fecha_alta	New_customer_Index	Seniority	Customer_Type_1st_month	...
\					
0	2009-09-15	0	69	P	...
1	2009-09-15	0	69	P	...
2	2009-09-15	0	69	P	...
3	2009-09-15	0	69	P	...
4	2009-09-15	0	69	P	...

	Taxes	Credit_card	Securities	Home_account	Payroll	Pensions.1	Direct_deb
it \							
0	0	0	0	0	0.0	0.0	
0							
1	0	0	0	0	0.0	0.0	
0							
2	0	0	0	0	0.0	0.0	
0							
3	0	0	0	0	0.0	0.0	
0							
4	0	0	0	0	0.0	0.0	
0							

	log_gross_income	sqrt_gross_income	winsorized_gross_income
0	12.0109	405.632321	164537.58
1	12.0109	405.632321	164537.58
2	12.0109	405.632321	164537.58
3	12.0109	405.632321	164537.58
4	12.0109	405.632321	164537.58

[5 rows x 48 columns]

Basic Information:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 11961703 entries, 0 to 11961702

Data columns (total 48 columns):

#	Column	Dtype
0	fecha_dato	object
1	Customer_Code	int64
2	Employee_index	object
3	Country_of_Residence	object
4	Sex	object
5	Age	int64
6	fecha_alta	object
7	New_customer_Index	int64
8	Seniority	int64
9	Customer_Type_1st_month	object
10	Customer_relation_Type_1st_month	object
11	Residence_index	object
12	Foreigner_index	object
13	conyuemp	object
14	Channel_used_to_join	object
15	indfall	object
16	cod_prov	float64
17	Province_name	object
18	Activity_index	float64
19	Gross_income	float64
20	Segmentation	object

```

21 Savings_account          int64
22 Guarantees                int64
23 Current_accounts          int64
24 Derivada_account          int64
25 Payroll_account           int64
26 Junior_account            int64
27 Mas_Particular_account    int64
28 Particular_account        int64
29 Particular_plus_account   int64
30 Short_term_deposits       int64
31 Medium_term_deposits      int64
32 Long_term_deposits        int64
33 e-accounts                int64
34 Funds                     int64
35 Mortgage                  int64
36 Pensions                  int64
37 Loans                     int64
38 Taxes                     int64
39 Credit_card               int64
40 Securities                 int64
41 Home_account              int64
42 Payroll                   float64
43 Pensions.1                 float64
44 Direct_debit              int64
45 log_gross_income           float64
46 sqrt_gross_income          float64
47 winsorized_gross_income    float64

```

dtypes: float64(8), int64(26), object(14)

memory usage: 4.3+ GB

None

Summary Statistics:

	Customer_Code	Age	New_customer_Index	Seniority \
count	1.196170e+07	1.196170e+07	1.196170e+07	1.196170e+07
mean	8.420010e+05	3.877639e+01	5.120843e-02	7.829003e+01
std	4.265184e+05	1.540047e+01	2.204226e-01	6.574827e+01
min	1.589000e+04	5.000000e+00	0.000000e+00	0.000000e+00
25%	4.656960e+05	2.400000e+01	0.000000e+00	2.300000e+01
50%	9.401310e+05	3.800000e+01	0.000000e+00	4.900000e+01
75%	1.200944e+06	4.900000e+01	0.000000e+00	1.330000e+02
max	1.548217e+06	9.100000e+01	1.000000e+00	2.560000e+02

	cod_prov	Activity_index	Gross_income	Savings_account \
count	1.196170e+07	1.196170e+07	1.196170e+07	1.196170e+07
mean	2.661878e+01	4.592029e-01	8.738710e+04	1.001530e-04
std	1.294053e+01	4.983328e-01	6.695058e+04	1.000714e-02
min	1.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
25%	1.500000e+01	0.000000e+00	4.106738e+04	0.000000e+00
50%	2.800000e+01	0.000000e+00	8.080068e+04	0.000000e+00
75%	3.500000e+01	1.000000e+00	1.274842e+05	0.000000e+00
max	5.200000e+01	1.000000e+00	2.801630e+05	1.000000e+00

	Guarantees	Current_accounts	...	Taxes	Credit_card \
count	1.196170e+07	1.196170e+07	...	1.196170e+07	1.196170e+07
mean	2.031483e-05	6.638349e-01	...	5.211984e-02	4.591127e-02
std	4.507152e-03	4.723962e-01	...	2.222687e-01	2.092927e-01
min	0.000000e+00	0.000000e+00	...	0.000000e+00	0.000000e+00

25%	0.000000e+00	0.000000e+00	...	0.000000e+00	0.000000e+00
50%	0.000000e+00	1.000000e+00	...	0.000000e+00	0.000000e+00
75%	0.000000e+00	1.000000e+00	...	0.000000e+00	0.000000e+00
max	1.000000e+00	1.000000e+00	...	1.000000e+00	1.000000e+00

	Securities	Home_account	Payroll	Pensions.1	Direct_debi
t \					
count	1.196170e+07	1.196170e+07	1.196170e+07	1.196170e+07	1.196170e+07
mean	2.494210e-02	3.895014e-03	5.788791e-02	6.198716e-02	1.333890e-01
std	1.559487e-01	6.228839e-02	2.335314e-01	2.411322e-01	3.399947e-01
min	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
25%	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
50%	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
75%	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
max	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00

	log_gross_income	sqrt_gross_income	winsorized_gross_income
count	1.196170e+07	1.196170e+07	1.196170e+07
mean	9.133037e+00	2.553196e+02	8.738710e+04
std	4.646379e+00	1.489934e+02	6.695058e+04
min	0.000000e+00	0.000000e+00	0.000000e+00
25%	1.062299e+01	2.026509e+02	4.106738e+04
50%	1.129975e+01	2.842546e+02	8.080068e+04
75%	1.175576e+01	3.570494e+02	1.274842e+05
max	1.254313e+01	5.293043e+02	2.801630e+05

[8 rows x 34 columns]

Missing Values:

fecha_dato	0
Customer_Code	0
Employee_index	0
Country_of_Residence	0
Sex	0
Age	0
fecha_alta	0
New_customer_Index	0
Seniority	0
Customer_Type_1st_month	0
Customer_relation_Type_1st_month	0
Residence_index	0
Foreigner_index	0
conyuemp	0
Channel_used_to_join	0
indfall	0
cod_prov	0
Province_name	0
Activity_index	0
Gross_income	0

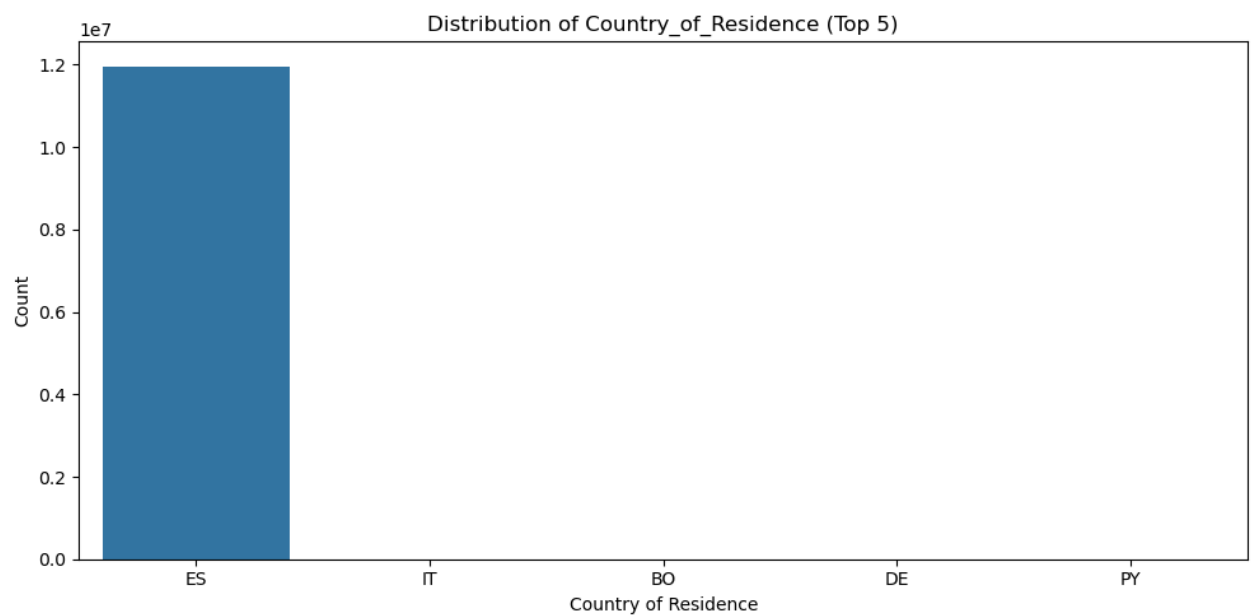
Segmentation	0
Savings_account	0
Guarantees	0
Current_accounts	0
Derivada_account	0
Payroll_account	0
Junior_account	0
Mas_Particular_account	0
Particular_account	0
Particular_plus_account	0
Short_term_deposits	0
Medium_term_deposits	0
Long_term_deposits	0
e-accounts	0
Funds	0
Mortgage	0
Pensions	0
Loans	0
Taxes	0
Credit_card	0
Securities	0
Home_account	0
Payroll	0
Pensions.1	0
Direct_debit	0
log_gross_income	0
sqrt_gross_income	0
winsorized_gross_income	0
dtype:	int64

```
In [25]: # Assuming df_train is your DataFrame
top_5_countries = df_train['Country_of_Residence'].value_counts().nlargest(5)

# Filter the DataFrame to include only the top 5 countries
df_top_5 = df_train[df_train['Country_of_Residence'].isin(top_5_countries)]

# Plot the distribution
plt.figure(figsize=(10, 5))
sns.countplot(x='Country_of_Residence', data=df_top_5, order=top_5_countries.index)
plt.title('Distribution of Country_of_Residence (Top 5)')
plt.xlabel('Country of Residence')
plt.ylabel('Count')
plt.tight_layout()
plt.show()

print(df_train['Country_of_Residence'].value_counts())
```



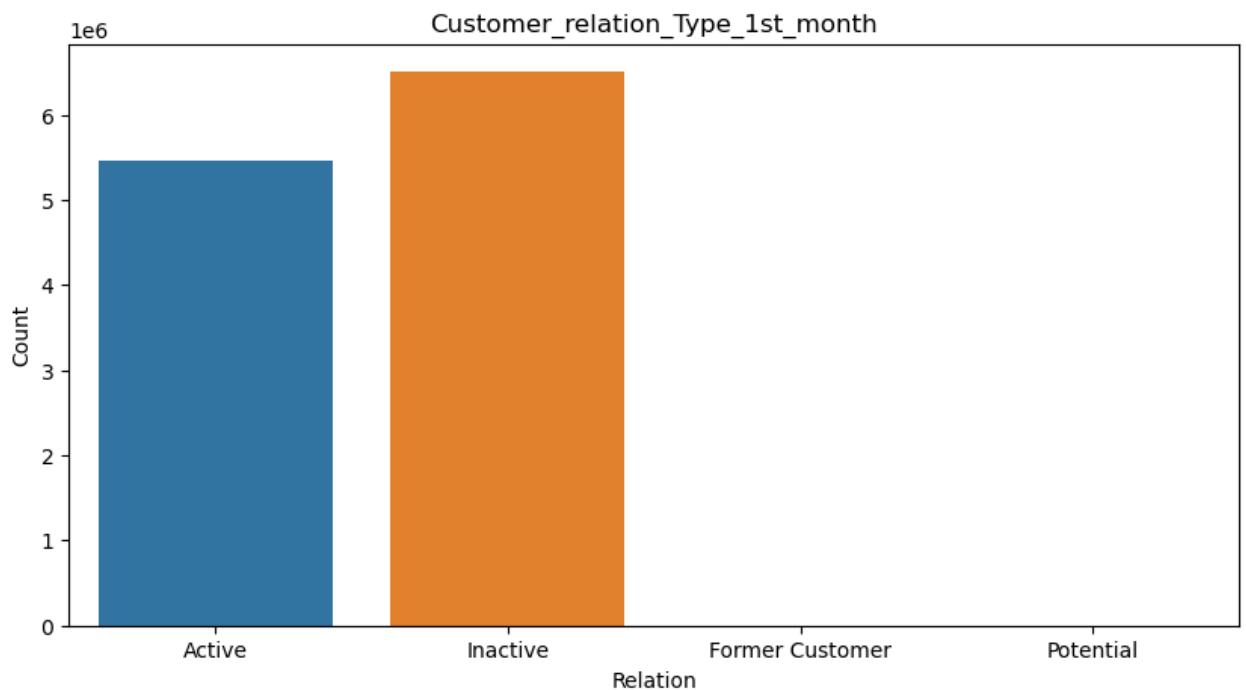
```
Country_of_Residence
ES      11961636
IT         17
BO         17
DE         17
PY         16
Name: count, dtype: int64
```

```
In [26]: # Check if the mapping was successful
print(df_train['Customer_relation_Type_1st_month'].value_counts())

# Plot the distribution for the "Sex" column
plt.figure(figsize=(10, 5))
ax = sns.countplot(x='Customer_relation_Type_1st_month', data=df_train, o
plt.title('Customer_relation_Type_1st_month')
plt.xlabel('Relation')
plt.ylabel('Count')
```

```
Customer_relation_Type_1st_month
Inactive      6504875
Active        5456716
Former Customer    112
Name: count, dtype: int64
```

```
Out[26]: Text(0, 0.5, 'Count')
```




```
In [27]: # Assuming df_train is your DataFrame
top_10_countries = df_train['Province_name'].value_counts().nlargest(10).

# Filter the DataFrame to include only the top 5 countries
df_top_10 = df_train[df_train['Province_name'].isin(top_10_countries)]

# Plot the distribution
plt.figure(figsize=(10, 5))
sns.countplot(x='Province_name', data=df_top_10, order=top_10_countries)
plt.title('Distribution of Province_name (Top 10)')
plt.xlabel('Province name')
plt.ylabel('Count')
plt.tight_layout()
plt.show()

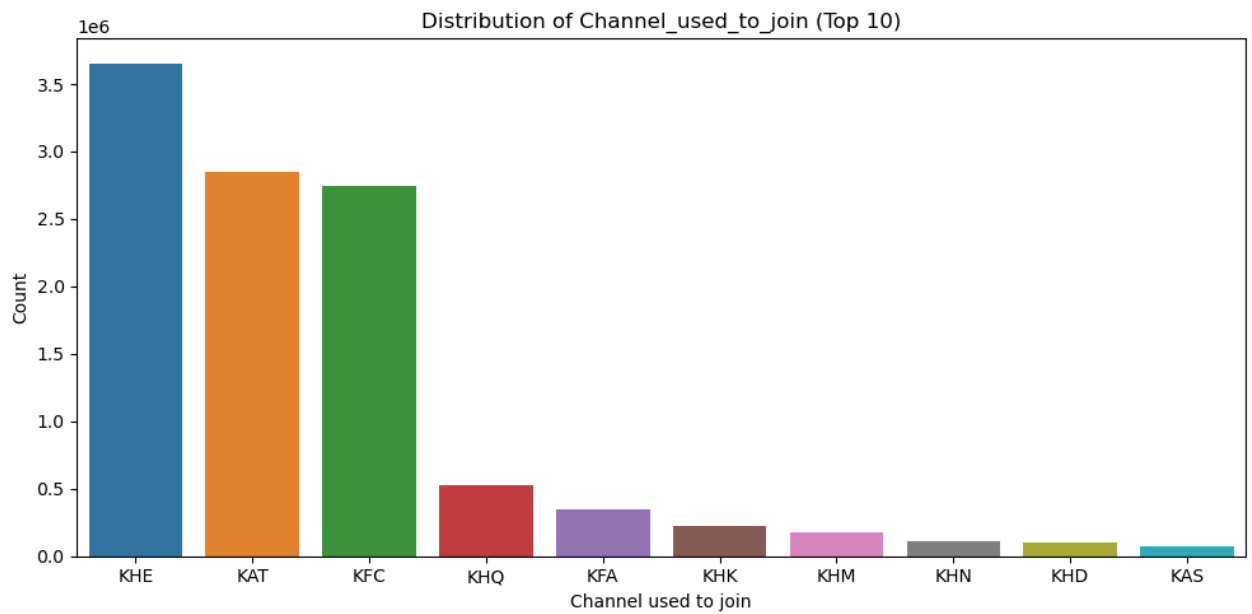
#print(df_train['Province_name'].value_counts())
```



```
In [28]: # Assuming df_train is your DataFrame
top_10_countries = df_train['Channel_used_to_join'].value_counts().nlarge
```

```
# Filter the DataFrame to include only the top 5 countries
df_top_10 = df_train[df_train['Channel_used_to_join'].isin(top_10_countries)]

# Plot the distribution
plt.figure(figsize=(10, 5))
sns.countplot(x='Channel_used_to_join', data=df_top_10, order=top_10_countries)
plt.title('Distribution of Channel_used_to_join (Top 10)')
plt.xlabel('Channel used to join')
plt.ylabel('Count')
plt.tight_layout()
plt.show()
```



```
In [29]: # Assuming df_cleaned is your DataFrame containing the relevant data

# Identify the top 5 channels
top_5_channels = df_train['Channel_used_to_join'].value_counts().nlargest(5)

# Filter the DataFrame to include only the top 5 channels
df_top_5_channels = df_train[df_train['Channel_used_to_join'].isin(top_5_channels)]

# Identify the top 10 provinces
top_10_provinces = df_train['Province_name'].value_counts().nlargest(10)

# Filter the DataFrame to include only the top 10 provinces
df_top_10_provinces = df_train[df_train['Province_name'].isin(top_10_provinces)]

# Create subplots to compare age distributions
fig, axes = plt.subplots(2, 1, figsize=(14, 14))

# Plot for Top 5 Channels
sns.boxplot(x='Channel_used_to_join', y='Age', data=df_top_5_channels, ax=axes[0])
axes[0].set_title('Box Plot of Age by Top 5 Channels Used to Join')
axes[0].set_xlabel('Channel Used to Join')
axes[0].set_ylabel('Age')

# Plot for Top 10 Provinces
sns.boxplot(x='Province_name', y='Age', data=df_top_10_provinces, ax=axes[1])
axes[1].set_title('Box Plot of Age by Top 10 Provinces')
```


```

axes[1].set_xlabel('Province Name')
axes[1].set_ylabel('Age')

# Rotate x-axis labels for better readability
for ax in axes:
    for label in ax.get_xticklabels():
        label.set_rotation(45)

plt.tight_layout()
plt.show()


```



```

In [31]: # Check if 'Segmentation' and 'Seniority' columns exist
if 'Segmentation' in df_train.columns and 'Seniority' in df_train.columns:
    plt.figure(figsize=(12, 8))
    sns.boxplot(x='Segmentation', y='Seniority', data=df_train)
    plt.title('Box Plot of Seniority by Segmentation')
    plt.xlabel('Segmentation')
    plt.ylabel('Seniority')
    plt.xticks(rotation=45) # Rotate x-axis labels if necessary
    plt.show()
else:
    print("One or more columns do not exist in the dataset.")

```



```
In [ ]: ## Visualize the distribution of numerical features
# numerical_features = df_train.select_dtypes(include=['int64', 'float64'])
## Box plots to visualize the spread and outliers in numerical features
# plt.figure(figsize=(14, 7))
# for i, col in enumerate(numerical_features):
#     plt.subplot(len(numerical_features) // 2 + 1, 2, i + 1)
#     sns.boxplot(y=col, data=df_train)
#     plt.title(f'Box Plot of {col}')
# plt.tight_layout()
# plt.show()
```

```
In [48]: # Identify the top 5 channels
top_5_channels = df_train['Channel_used_to_join'].value_counts().nlargest(5)

# Filter the DataFrame to include only the top 5 channels
df_top_5_channels = df_train[df_train['Channel_used_to_join'].isin(top_5_channels)]

# Create a pivot table for the mean age by segmentation and top 5 channel
pivot_table = df_top_5_channels.pivot_table(values='Age', index='Segmentation')

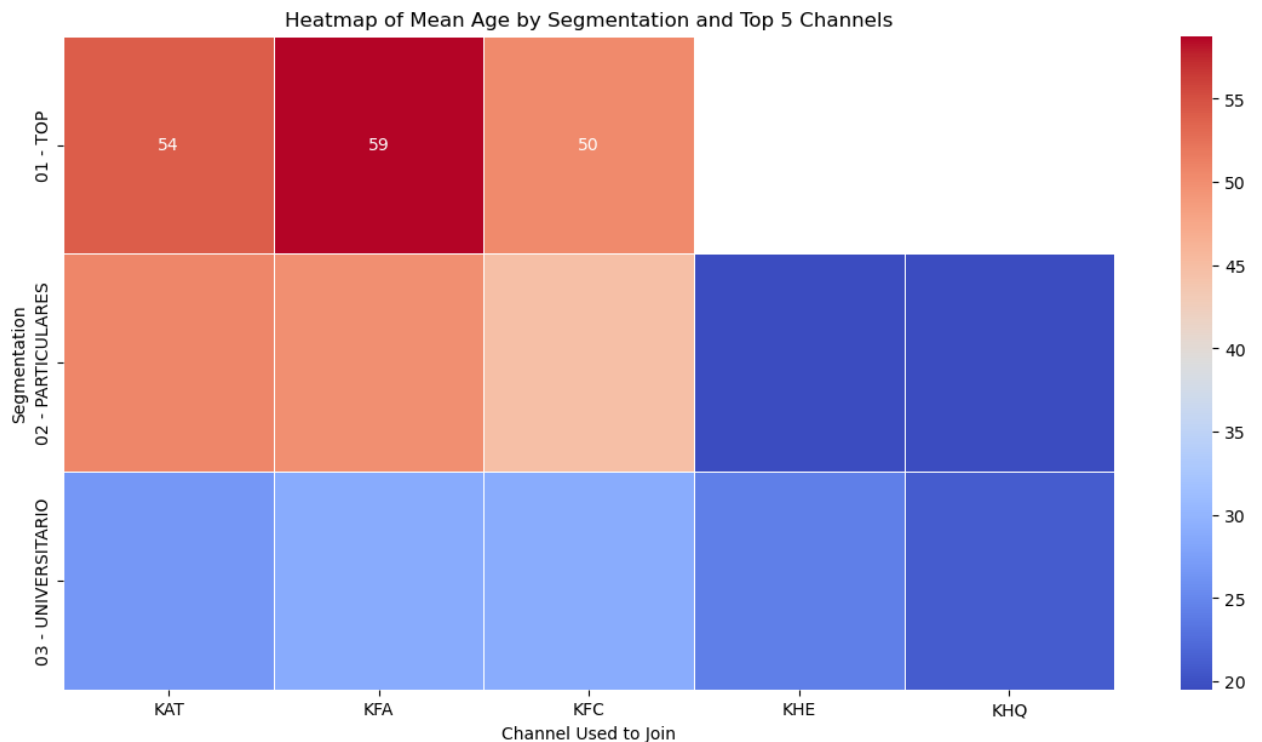
# Print the pivot table for verification
print("\nPivot Table of Mean Age by Segmentation and Top 5 Channels:")
print(pivot_table)

# Create the heatmap
plt.figure(figsize=(14, 7))
sns.heatmap(pivot_table, annot=True, cmap='coolwarm', linewidths=.5)
plt.title('Heatmap of Mean Age by Segmentation and Top 5 Channels')
```

```
plt.xlabel('Channel Used to Join')
plt.ylabel('Segmentation')
plt.show()
```

Pivot Table of Mean Age by Segmentation and Top 5 Channels:

Channel_used_to_join	KAT	KFA	KFC	KHE	KH
Q					
Segmentation					
01 - TOP	54.013806	58.764498	50.332913	NaN	Na
N					
02 - PARTICULARES	50.785190	49.892479	44.760641	19.601375	19.51855
7					
03 - UNIVERSITARIO	26.712131	28.740277	28.970408	24.147726	21.02599
6					

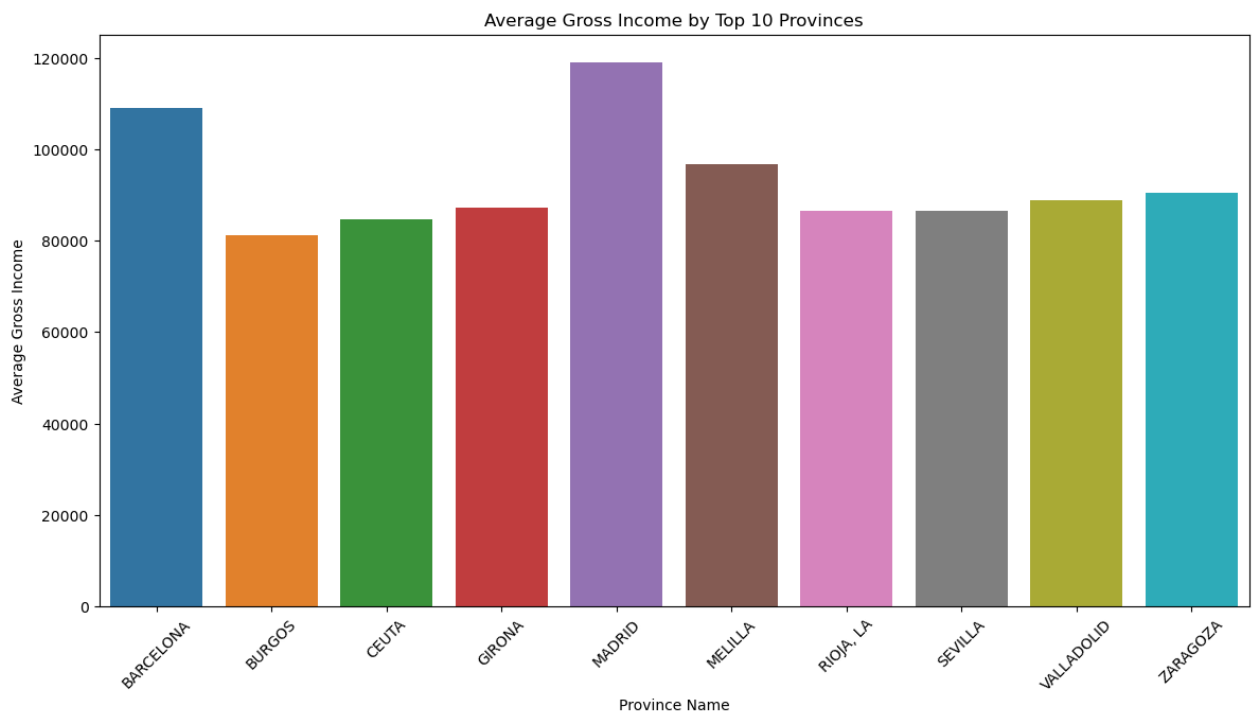


```
In [68]: # Calculate the top 10 provinces by average Gross_income
top_10_provinces = df_train.groupby('Province_name')['Gross_income'].mean
print("Top 10 Provinces by Average Gross Income:")
print(top_10_provinces)
# Filter the DataFrame to include only the top 10 provinces
df_top_10_provinces = df_train[df_train['Province_name'].isin(top_10_prov

# Bar plot of average Gross_income by top 10 Province_name
plt.figure(figsize=(14, 7))
avg_income_by_top_10_provinces = df_top_10_provinces.groupby('Province_na
sns.barplot(x='Province_name', y='Gross_income', data=avg_income_by_top_1
plt.title('Average Gross Income by Top 10 Provinces')
plt.xlabel('Province Name')
plt.ylabel('Average Gross Income')
plt.xticks(rotation=45)
plt.show()
```

Top 10 Provinces by Average Gross Income:

```
Index(['MADRID', 'BARCELONA', 'MELILLA', 'ZARAGOZA', 'VALLADOLID', 'GIRON
A',
      'SEVILLA', 'RIOJA, LA', 'CEUTA', 'BURGOS'],
      dtype='object', name='Province_name')
```

```
In [93]: # Calculate the top 5 residence_index values by their frequency
top_5_residence_index = df_train['Residence_index'].value_counts().nlarge

# Calculate the top 5 Channel_used_to_join values by their frequency
top_5_channels = df_train['Channel_used_to_join'].value_counts().nlargest

print("Top 5 Residence Index Values:")
print(top_5_residence_index)
print("\nTop 5 Channels Used to Join:")
print(top_5_channels)

# Filter the DataFrame to include only the top 5 residence_index and top
df_top_5 = df_train[df_train['Residence_index'].isin(top_5_residence_index)]

# Calculate the counts for each combination of residence_index and Channel
residence_channel_counts_top_5 = df_top_5.groupby(['Residence_index', 'Channel_used_to_join']).count()

# Plot the counts using a bar plot
plt.figure(figsize=(14, 7))
residence_channel_counts_top_5.plot(kind='bar', stacked=True, colormap='viridis')
plt.title('Counts of Top 5 Residence Index by Top 5 Channels Used to Join')
plt.xlabel('Residence Index')
plt.ylabel('Count')
plt.legend(title='Channel Used to Join', bbox_to_anchor=(1.05, 1), loc='upper right')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

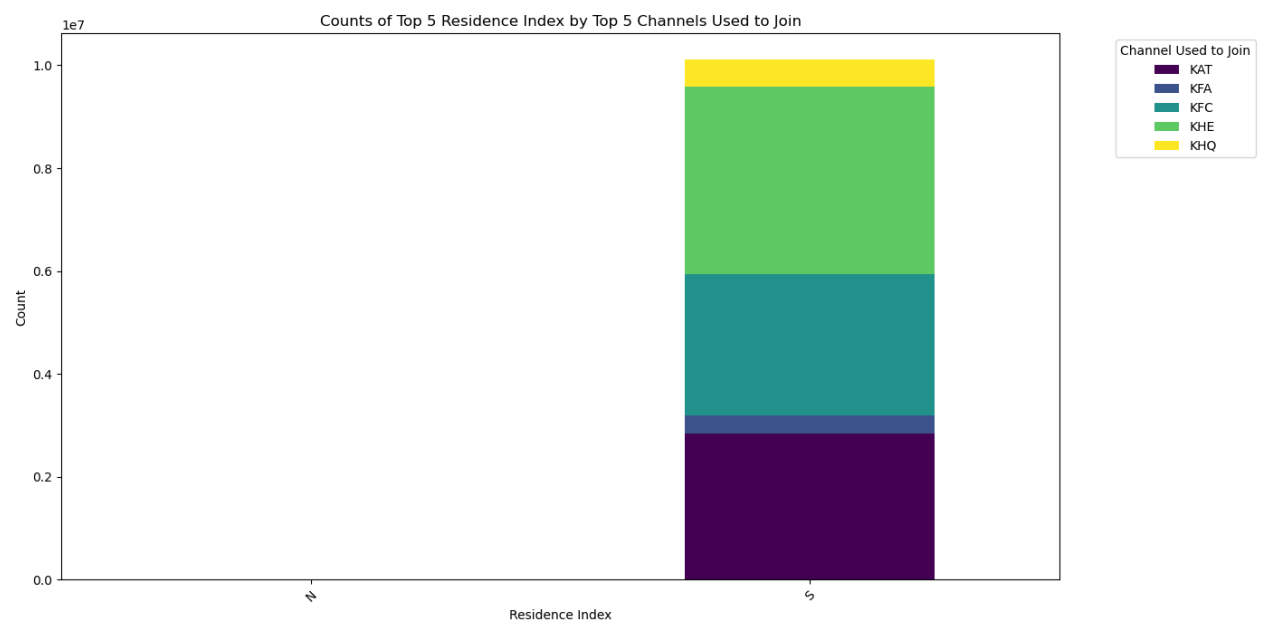
Top 5 Residence Index Values:

Index(['S', 'N'], dtype='object', name='Residence_index')

Top 5 Channels Used to Join:

Index(['KHE', 'KAT', 'KFC', 'KHQ', 'KFA'], dtype='object', name='Channel_used_to_join')

<Figure size 1400x700 with 0 Axes>



```
In [96]: # Plot the counts using a heatmap
plt.figure(figsize=(14, 7))
sns.heatmap(residence_channel_counts_top_5, annot=True, fmt='d', cmap='YlGnBu')
plt.title('Heatmap of Top 5 Residence Index by Top 5 Channels Used to Join')
plt.xlabel('Channel Used to Join')
plt.ylabel('Residence Index')
plt.xticks(rotation=45)
plt.show()
```

