

```
In [3]: import pandas as pd
import numpy as np
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
import matplotlib.pyplot as plt
```

```
In [5]: # Read the CSV file
df_train = pd.read_csv("train_new.csv")
```

```
In [9]: df_train['log_gross_income'] = np.log(df_train['Gross_income'] + 1)
```

```
In [13]: df_train['sqrt_gross_income'] = np.sqrt(df_train['Gross_income'])
```

```
In [16]: # Display the updated DataFrame
print(df_train.head())
```

	fecha_dato	Customer_Code	Employee_index	Country_of_Residence	Sex	A
ge \						
0	2015-01-28	1375586	Not employed	ES	Men	
35						
1	2015-01-28	1050611	Not employed	ES	Women	
23						
2	2015-01-28	1050612	Not employed	ES	Women	
23						
3	2015-01-28	1050613	Not employed	ES	Men	
22						
4	2015-01-28	1050614	Not employed	ES	Women	
23						

	fecha_alta	New_customer_Index	Seniority	Customer_Type_1st_month	...
\					
0	2015-01-12	0	6	P	...
1	2012-08-10	0	35	P	...
2	2012-08-10	0	35	P	...
3	2012-08-10	0	35	P	...
4	2012-08-10	0	35	P	...

	Loans	Taxes	Credit_card	Securities	Home_account	Payroll	Pensions.1	\
0	0	0	0	0	0	0.0	0.0	
1	0	0	0	0	0	0.0	0.0	
2	0	0	0	0	0	0.0	0.0	
3	0	0	0	0	0	0.0	0.0	
4	0	0	0	0	0	0.0	0.0	

	Direct_debit	log_gross_income	sqrt_gross_income
0	0	11.376179	295.327107
1	0	10.478688	188.543735
2	0	11.713252	349.541285
3	0	11.693383	346.086030
4	0	0.000000	0.000000

[5 rows x 47 columns]

```
In [21]: from scipy.stats.mstats import winsorize
```

```
df_train['winsorized_gross_income'] = winsorize(df_train['Gross_income'],
```

```
In [37]: # Calculate the first quartile (Q1) and the third quartile (Q3)
Q1 = df_train['Gross_income'].quantile(0.25)
Q3 = df_train['Gross_income'].quantile(0.75)

# Calculate the Interquartile Range (IQR)
IQR = Q3 - Q1

# Define the lower and upper thresholds
lower_threshold = Q1 - 1.5 * IQR
upper_threshold = Q3 + 1.5 * IQR

# Filter the DataFrame to remove outliers
df_cleaned = df_train[(df_train['Gross_income'] >= lower_threshold) & (df

# Display the result (optional)
print(df_cleaned)
```

	fecha_dato	Customer_Code	Employee_index	Country_of_Residence	\
0	2015-01-28	1375586	Not employed	ES	
1	2015-01-28	1050611	Not employed	ES	
2	2015-01-28	1050612	Not employed	ES	
3	2015-01-28	1050613	Not employed	ES	
4	2015-01-28	1050614	Not employed	ES	
...
13379651	2016-05-28	1166766	Not employed	ES	
13379652	2016-05-28	1166765	Not employed	ES	
13379653	2016-05-28	1166764	Not employed	ES	
13379654	2016-05-28	1166763	Not employed	ES	
13379655	2016-05-28	1166789	Not employed	ES	

	Sex	Age	fecha_alta	New_customer_Index	Seniority	\
0	Men	35	2015-01-12	0	6	
1	Women	23	2012-08-10	0	35	
2	Women	23	2012-08-10	0	35	
3	Men	22	2012-08-10	0	35	
4	Women	23	2012-08-10	0	35	
...
13379651	Women	25	2013-08-14	0	33	
13379652	Women	22	2013-08-14	0	33	
13379653	Women	23	2013-08-14	0	33	
13379654	Men	47	2013-08-14	0	33	
13379655	Men	22	2013-08-14	0	33	

	Customer_Type_1st_month	...	Taxes	Credit_card	Securities	\
0	P	...	0	0	0	
1	P	...	0	0	0	
2	P	...	0	0	0	
3	P	...	0	0	0	
4	P	...	0	0	0	
...
13379651	1	...	0	0	0	
13379652	1	...	0	0	0	
13379653	1	...	0	0	0	
13379654	1	...	0	0	0	
13379655	P	...	0	0	0	

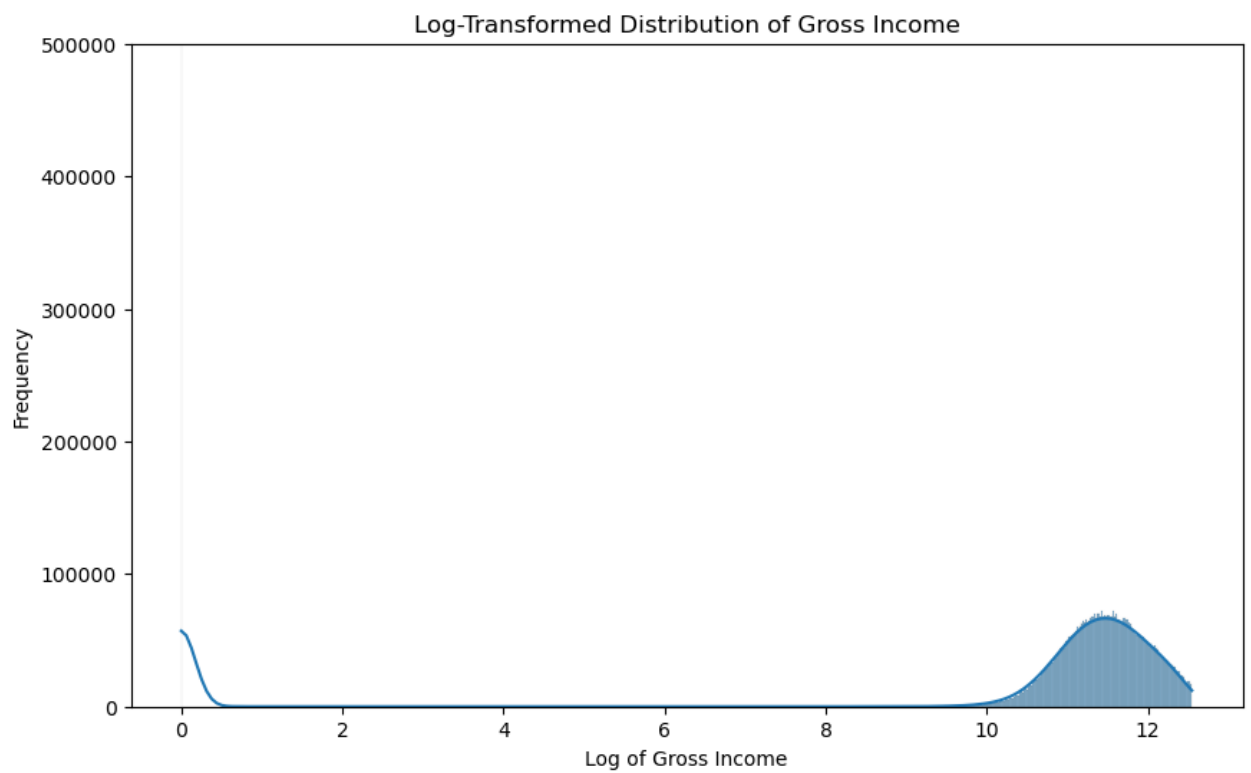
	Home_account	Payroll	Pensions.1	Direct_debit	log_gross_income	\
0	0	0.0	0.0	0	11.376179	
1	0	0.0	0.0	0	10.478688	
2	0	0.0	0.0	0	11.713252	
3	0	0.0	0.0	0	11.693383	
4	0	0.0	0.0	0	0.000000	
...
13379651	0	0.0	0.0	0	10.838526	
13379652	0	0.0	0.0	0	10.689970	
13379653	0	0.0	0.0	0	10.057752	
13379654	0	0.0	0.0	0	0.000000	
13379655	0	0.0	0.0	0	12.204040	

	sqrt_gross_income	winsorized_gross_income
0	295.327107	87218.10
1	188.543735	35548.74
2	349.541285	122179.11
3	346.086030	119775.54
4	0.000000	0.00
...
13379651	225.710545	50945.25
13379652	209.552309	43912.17
13379653	152.757946	23334.99
13379654	0.000000	0.00
13379655	446.758122	199592.82

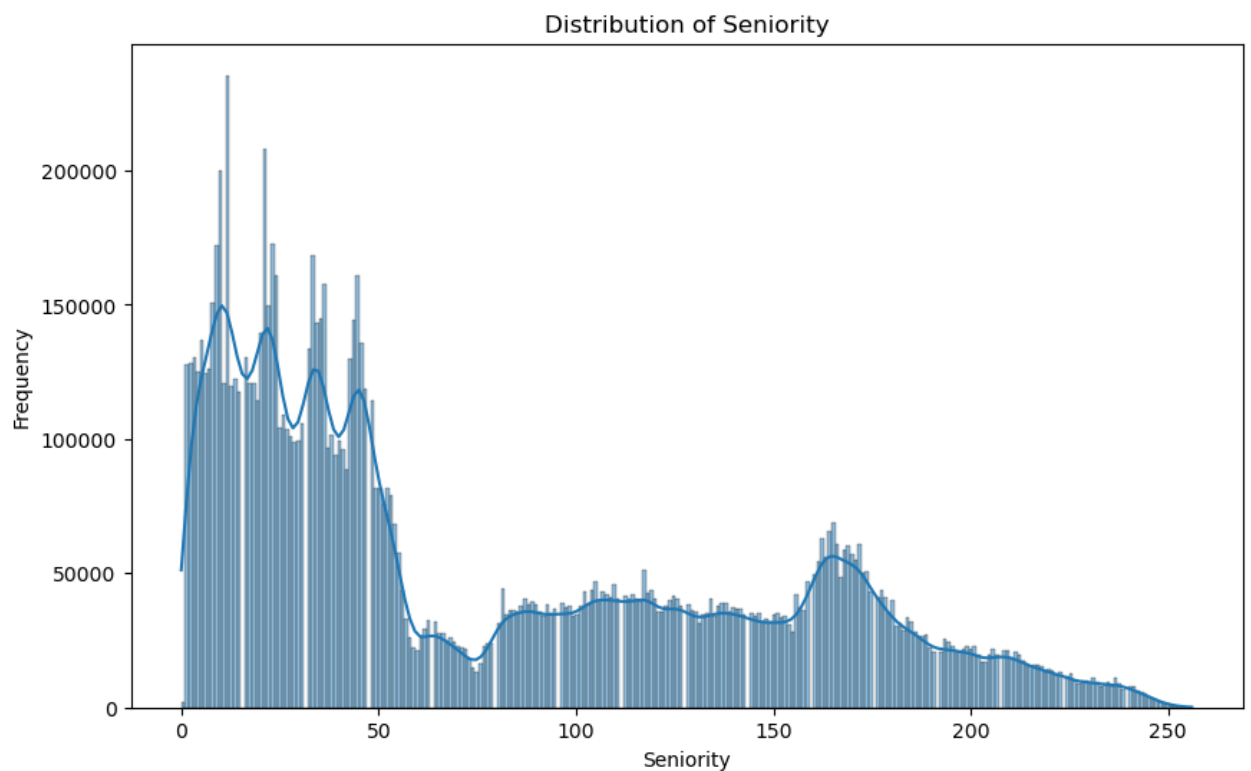
[12686208 rows x 48 columns]

```
In [46]: # Check if 'Gross_income' column exists
if 'Gross_income' in df_cleaned.columns:
    # Apply logarithmic transformation to the 'Gross_income' column
    log_data = np.log1p(df_cleaned['Gross_income'])

    # Create a histogram for the log-transformed 'Gross_income' column
    plt.figure(figsize=(10, 6))
    sns.histplot(log_data, kde=True)
    plt.title('Log-Transformed Distribution of Gross Income')
    plt.xlabel('Log of Gross Income')
    plt.ylabel('Frequency')
    plt.ylim(0, 500000) # Set the y-axis limit
    plt.show()
else:
    print("The 'Gross_income' column does not exist in the dataset.")
```



```
In [52]: # Check if Seniority column exists
if 'Seniority' in df_cleaned.columns:
    # Create a histogram for the Seniority column
    plt.figure(figsize=(10, 6))
    sns.histplot(df_cleaned['Seniority'], kde=True)
    plt.title('Distribution of Seniority')
    plt.xlabel('Seniority')
    plt.ylabel('Frequency')
    plt.show()
else:
    print("The Seniority column does not exist in the dataset.")
```



```
In [77]: # Calculate the first quartile (Q1) and the third quartile (Q3) for each
Q1 = df_train.groupby('Sex')['Age'].quantile(0.25)
Q3 = df_train.groupby('Sex')['Age'].quantile(0.75)

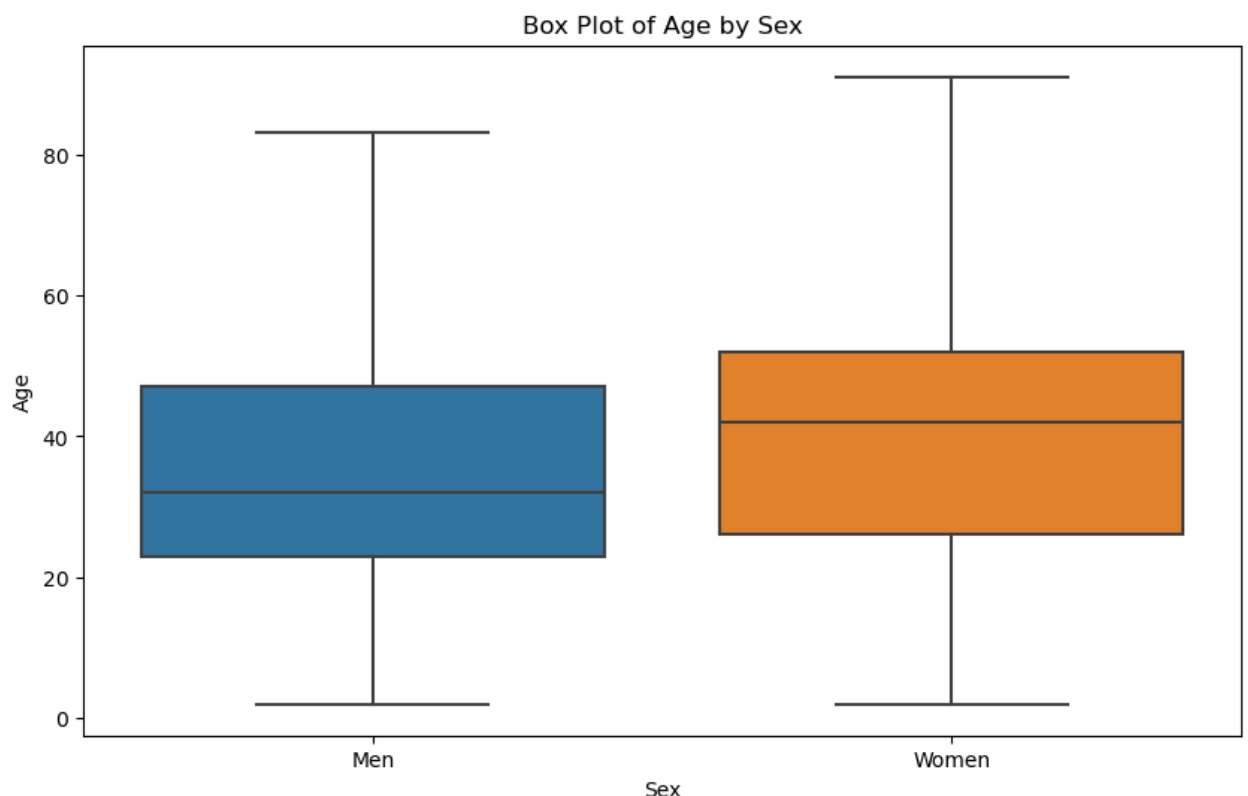
# Calculate the Interquartile Range (IQR) for each group
IQR = Q3 - Q1

# Define the lower and upper thresholds for each group
lower_threshold = Q1 - 1.5 * IQR
upper_threshold = Q3 + 1.5 * IQR

# Filter the DataFrame to remove outliers
def filter_outliers(group):
    lower = lower_threshold[group.name]
    upper = upper_threshold[group.name]
    return group[(group['Age'] >= lower) & (group['Age'] <= upper)]

df_cleaned = df_train.groupby('Sex').apply(filter_outliers).reset_index()

# Plot the cleaned data
plt.figure(figsize=(10, 6))
sns.boxplot(x='Sex', y='Age', data=df_cleaned)
plt.title('Box Plot of Age by Sex')
plt.xlabel('Sex')
plt.ylabel('Age')
plt.show()
```



```
In [87]: # Calculate the first quartile (Q1) and the third quartile (Q3) for each
Q1 = df_cleaned.groupby('Channel_used_to_join')['Age'].quantile(0.25)
Q3 = df_cleaned.groupby('Channel_used_to_join')['Age'].quantile(0.75)

# Calculate the Interquartile Range (IQR) for each channel
IQR = Q3 - Q1
```

```

# Define the lower and upper thresholds for each channel
lower_threshold = Q1 - 1.5 * IQR
upper_threshold = Q3 + 1.5 * IQR

# Function to filter out outliers based on IQR
def filter_outliers(group):
    lower = lower_threshold[group.name]
    upper = upper_threshold[group.name]
    return group[(group['Age'] >= lower) & (group['Age'] <= upper)]

# Apply the filter function to each group
df_no_outliers = df_cleaned.groupby('Channel_used_to_join').apply(filter_

```

	fecha_dato	Customer_Code	Employee_index	Country_of_Residence	\
0	2015-01-28	851959	Not employed	ES	
1	2015-02-28	851959	Not employed	ES	
2	2015-03-28	851959	Not employed	ES	
3	2015-04-28	851959	Not employed	ES	
4	2015-05-28	851959	Not employed	ES	
...
12617096	2016-05-28	1173589	Not employed	ES	
12617097	2016-05-28	1172121	Not employed	ES	
12617098	2016-05-28	1171673	Not employed	ES	
12617099	2016-05-28	1165203	Not employed	ES	
12617100	2016-05-28	1168589	Not employed	ES	

	Sex	Age	fecha_alta	New_customer_Index	Seniority	\
0	Men	36	2009-09-15	0	69	
1	Men	36	2009-09-15	0	69	
2	Men	36	2009-09-15	0	69	
3	Men	36	2009-09-15	0	69	
4	Men	36	2009-09-15	0	69	
...
12617096	Women	38	2014-05-13	0	26	
12617097	Women	32	2013-09-03	0	28	
12617098	Women	21	2013-09-02	0	12	
12617099	Women	24	2013-08-13	0	23	
12617100	Women	59	2013-09-17	0	23	

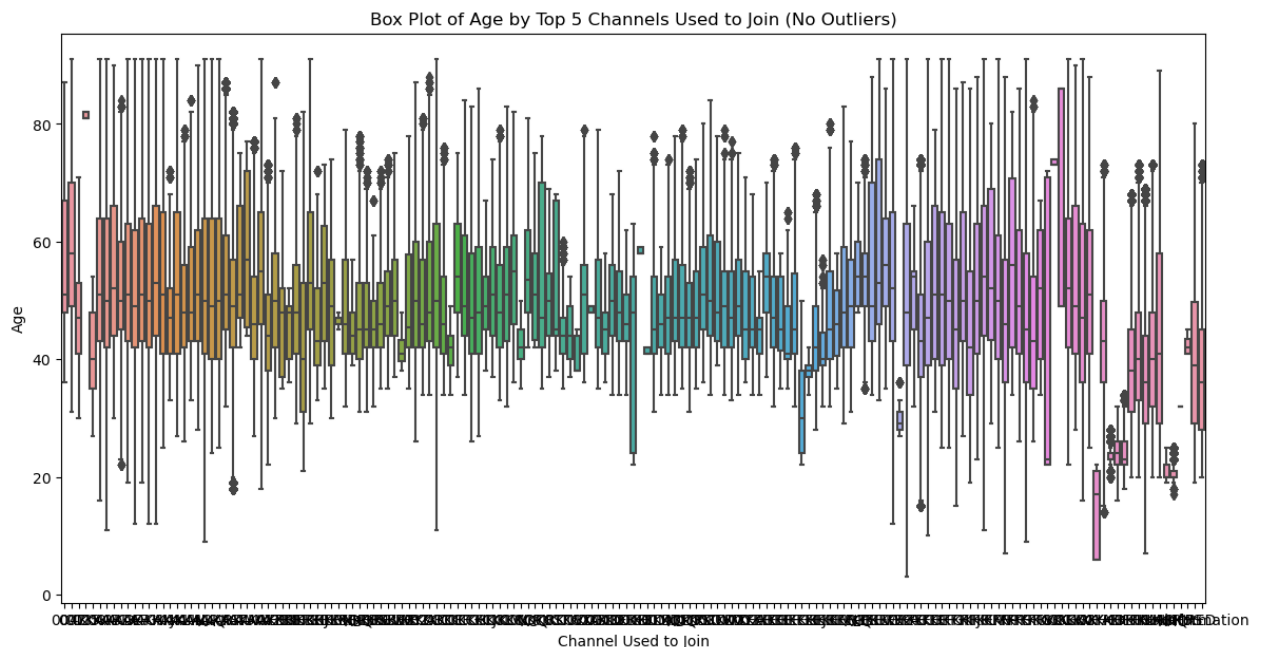
	Customer_Type_1st_month	...	Taxes	Credit_card	Securities	\
0	P	...	0	0	0	
1	P	...	0	0	0	
2	P	...	0	0	0	
3	P	...	0	0	0	
4	P	...	0	0	0	
...
12617096	1	...	0	0	0	
12617097	1	...	0	0	0	
12617098	P	...	0	0	0	
12617099	1	...	0	0	0	
12617100	P	...	0	0	0	

	Home_account	Payroll	Pensions.1	Direct_debit	log_gross_income	\
0	0	0.0	0.0	0	12.010900	
1	0	0.0	0.0	0	12.010900	
2	0	0.0	0.0	0	12.010900	
3	0	0.0	0.0	0	12.010900	

4	0	0.0	0.0	0	12.010900
...
12617096	0	0.0	0.0	0	11.286214
12617097	0	0.0	0.0	0	10.746474
12617098	0	0.0	0.0	0	11.816476
12617099	0	0.0	0.0	0	11.253567
12617100	0	0.0	0.0	1	0.000000

	sqrt_gross_income	winsorized_gross_income
0	405.632321	164537.58
1	405.632321	164537.58
2	405.632321	164537.58
3	405.632321	164537.58
4	405.632321	164537.58
...
12617096	282.336873	79714.11
12617097	215.557185	46464.90
12617098	368.055784	135465.06
12617099	277.765387	77153.61
12617100	0.000000	0.00

[12617101 rows x 48 columns]

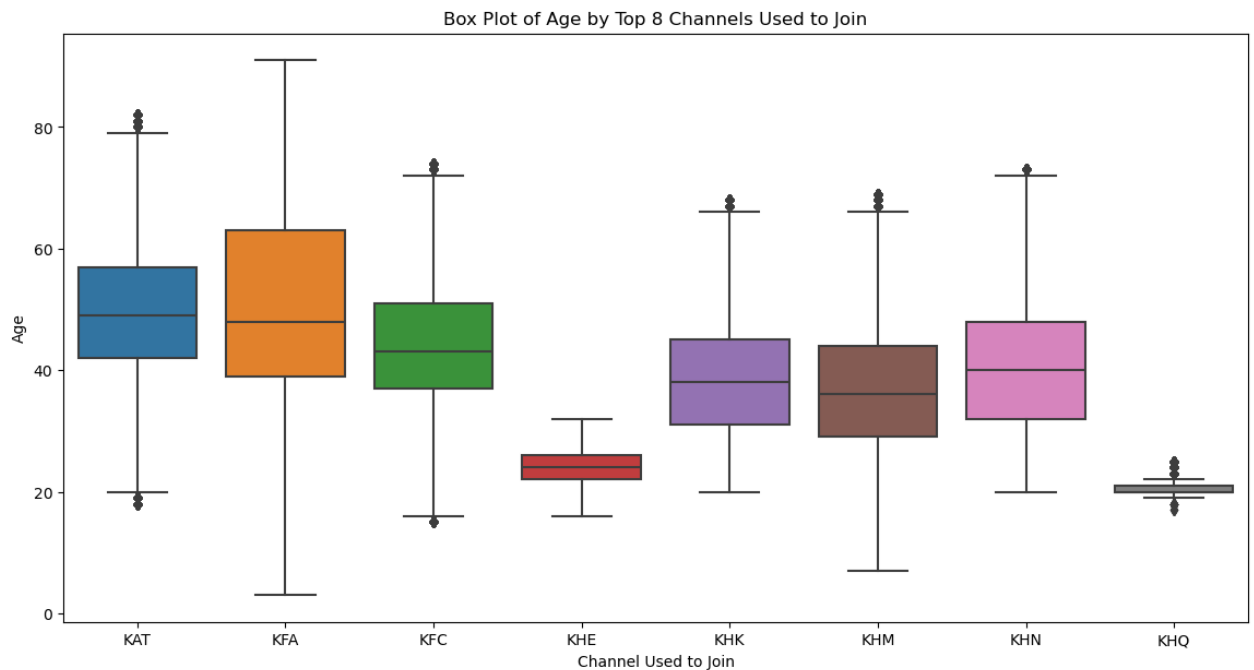


```
In [93]: # Calculate the count of each channel
channel_counts = df_no_outliers['Channel_used_to_join'].value_counts()

# Get the top 8 channels
top_8_channels = channel_counts.head(8).index

# Filter the DataFrame to include only the top 5 channels
df_top_8_channels = df_no_outliers[df_no_outliers['Channel_used_to_join']]

# Create the box plot
plt.figure(figsize=(14, 7))
sns.boxplot(x='Channel_used_to_join', y='Age', data=df_top_8_channels)
plt.title('Box Plot of Age by Top 8 Channels Used to Join')
plt.xlabel('Channel Used to Join')
plt.ylabel('Age')
plt.show()
```



```
In [99]: # List of columns to exclude from the correlation matrix
columns_to_exclude = ['Customer_Code'] # Add other columns to exclude if

# Drop the specified columns and select numerical columns
relevant_numerical_df = df_no_outliers.drop(columns=columns_to_exclude).s


# Compute the correlation matrix
corr = relevant_numerical_df.corr()

# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(corr, dtype=bool))

# Set up the matplotlib figure
plt.figure(figsize=(15, 10))

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, annot=True, fmt='.2f', cmap='coolwarm', vmin=
            square=True, linewidths=.5, cbar_kws={"shrink": .5})

plt.title('Correlation Matrix')
plt.show()
```

```
In [100... # Scatter plot between 'Age' and 'Gross_income'
if 'Age' in df_no_outliers.columns and 'Gross_income' in df_no_outliers.c
    plt.figure(figsize=(10, 6))
    sns.scatterplot(x='Age', y='Gross_income', data=df_no_outliers)
    plt.title('Scatter Plot of Age vs Gross Income')
    plt.xlabel('Age')
    plt.ylabel('Gross Income')
    plt.show()
else:
    print("One or more columns do not exist in the dataset.")
```

